

# Maps and Sorted Maps



# Map ADT

---

A MAP is an ADT to efficiently store and retrieve values based on a uniquely identifying **search key**.

It stores key-value pairs  $(k,v)$ , which we call **entries**.

Keys are unique/no repeats (they uniquely identify the value); a key is mapped to a value.

The main operations of a MAP are **searching**, **inserting**, and **deleting** items.

Examples: student records, user accounts, etc

Typical keys are username, user ID, etc.

Maps are also known as **associative arrays**.

**Dictionary** ADT is related, although it normally refers to a similar ADT that allows repeated keys.

# The MAP ADT methods:

---

`get(k)`: returns the value  $v$  associated to key  $k$ , if such entry exists; otherwise returns null.

`put(k, v)`: if  $M$  does not have an entry with key  $k$ , then adds  $(k, v)$  and returns null; otherwise it replaces with  $v$  the value of the entry with key equal to  $k$  and returns the old value.

`remove(k)`: removes from  $M$  the entry with key  $k$  and returns its value; if  $M$  has no such entry, the returns null.

`size()`: returns the number of entries in  $M$ .

`isEmpty()`: boolean indicating if  $M$  is empty.

`keySet()`, `values()`, `entrySet()` returns an iterable collection of keys, values, key-value entries (respectively) stored in  $M$ .

# MAP ADT: examples

---

Applications/examples:

- **University information system:**

key= student id

value= student record (name, address, course grades)

- A **domain name system (DNS)** maps a host name (key, e.g. [www.wiley.com](http://www.wiley.com)) to a IP address (value, e.g. 208.215.179.146)

- A **social media site** maps a username which is the key (usually nonnumeric) to the user info which is the value (typically tons of personal info)

# SORTED MAP ADT methods:

In addition to the MAP methods such as  
get(k); put(k, v); remove(k); size(); isEmpty();  
keySet(); values(); entrySet()

A SORTED MAP also provides:

firstEntry(), lastEntry(): returns the entry with smallest key,  
largest key (respectively), or null if the map is empty.

subMap(k1,k2): returns an iterable list with all the entries  
greater than or equal to k1, but strictly less than k2.

lowerEntry(k), higherEntry(k), floorEntry(k), ceilingEntry(k)  
return the entry with, respectively:  
the greatest key  $< k$ , the smallest key  $> k$ ,  
the greatest key  $\leq k$ , the smallest key  $\geq k$ .

# Implementing MAPs:

---

- Using an Unordered Sequence
- Using an Ordered Sequence
- Using Search Trees - binary search trees, AVL trees, red-black trees, (2,4)-trees  
(starts this lecture and continue on next ones)
- Using Hash Table  
(discussed at a later lecture)

## Implementing SORTED MAPs:

- Using an Ordered Sequence
- Using Search Trees - binary search trees, AVL trees, red-black trees, (2,4)-trees

# Implementing MAPs with an Unordered Sequence

---

- *unordered sequence*



- *get, remove and put takes  $O(n)$  time*
- *The "insert" part takes  $O(1)$  time, but we need first to search for key duplicate which takes  $O(n)$ .*

# Implementing a Map an Ordered Sequence

---

- *array-based ordered sequence* (assumes keys can be ordered)




- searching takes  $O(\log n)$  time (binary search)
- inserting and removing takes  $O(n)$  time
- application to look-up tables (frequent searches, rare insertions and removals)



# Binary Search

- narrow down the search range in stages
- “high-low” game
- Example: `get(7)`



# Pseudocode for Binary Search

Algorithm **BinarySearch**(S, k, low, high)

if  $\text{low} > \text{high}$  then

    return **NO\_SUCH\_KEY**

else      $\text{mid} \leftarrow (\text{low} + \text{high}) / 2$

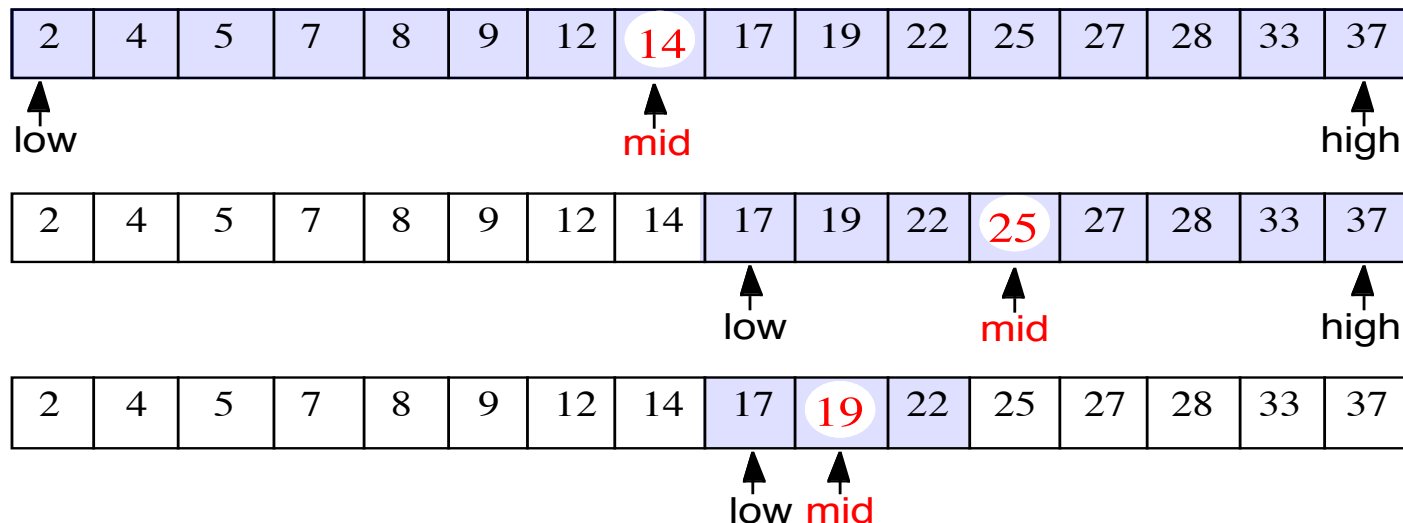
    if  $k = \text{key}(\text{mid})$  then

        return  $\text{key}(\text{mid})$

    else if  $k < \text{key}(\text{mid})$  then

        return **BinarySearch**(S, k, low, mid-1)

    else     return **BinarySearch**(S, k, mid+1, high)



# Running Time of Binary Search

---

- The range of candidate items to be searched is *halved after each comparison*

comparison	search range
0	$n$
1	$n/2$
2	$n/4$
...	...
$2^i$	$n/2^i$
$\log_2 n$	1

In the array-based implementation, access by rank takes  $O(1)$  time, thus *binary search runs in  $O(\log n)$  time*

# Binary Search Tree


---

Searching

Cost of Searching


Insertion


Deletion




# Binary Search Trees

- A binary search tree is a binary tree  $T$  such that
  - each internal node  $p$  stores an item  $(k, v)$  of a MAP.
  - keys stored at nodes in the left subtree of  $p$  are less than  $k$ .
  - keys stored at nodes in the right subtree of  $p$  are greater than  $k$ .
  - external nodes do not hold elements but serve as place holders (dummy leaves).







Question: How can we traverse the tree so that we visit the elements in increasing key order?

## IN-ORDER TRAVERSAL

always traverses the keys in increasing order  
in a binary search tree !!!



# MAP Operations using Binary Search Trees

---

Searching    `get(k):`  
                  `use TreeSearch(k)`

Inserting/updating value    `put(k, v):`  
                  `use TreeInsert(k,v)`

Removing    `remove(k):`  
                  `use TreeDelete(k)`

# Search

- To search for a key  $k$ , we trace a downward path starting at the root
- The next node visited depends on the outcome of the comparison of  $k$  with the key of the current node
- If we reach a leaf, the key is not found and we return this dummy leaf which will help with insertion if needed.
- Example: `TreeSearch(4)`

Algorithm `TreeSearch(k)`

`TreeSearch(root, k)`

Procedure `TreeSearch(p, k)`

if  $p$  is external then

return  $p$  {unsuccessful search}

else if  $k == \text{key}(p)$


return  $p$  {successful search}

else if  $k < \text{key}(p)$

return `TreeSearch(left(p), k)`

else {  $k > \text{key}(v)$  }


return `TreeSearch(right(p), k)`



# Search Example I

---

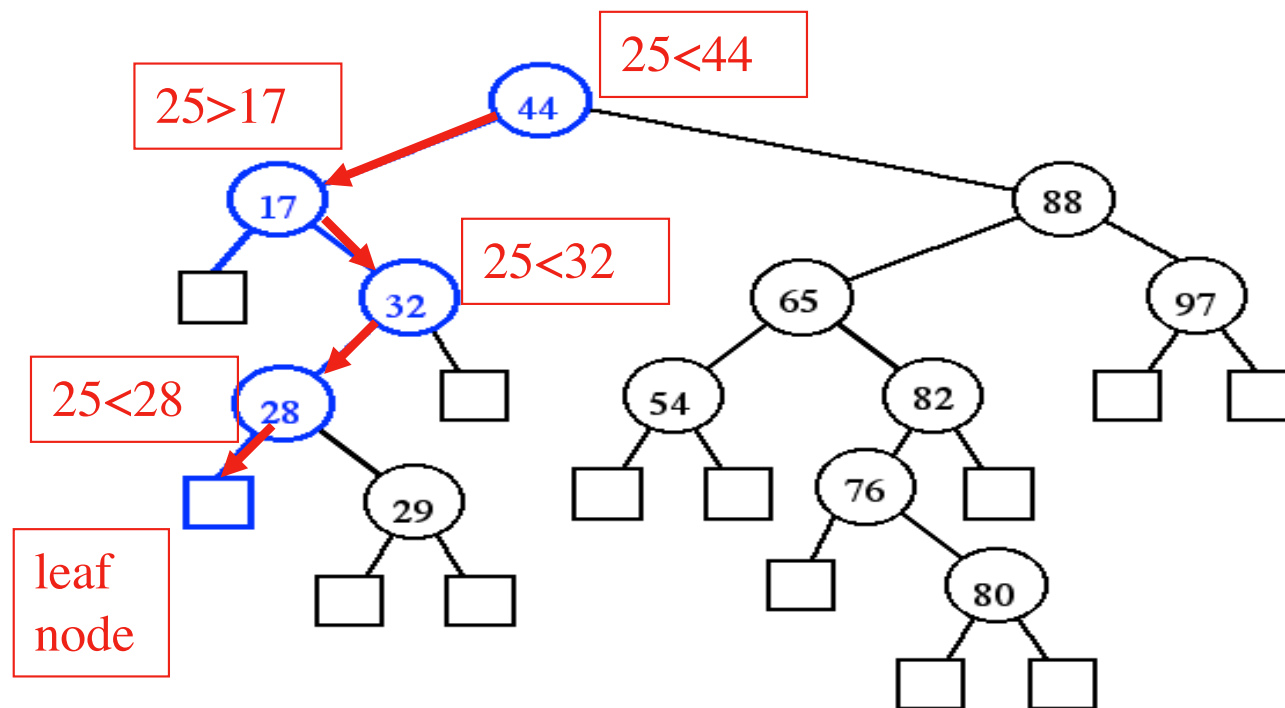
Successful **TreeSearch**(76)



- A successful search traverses a path starting at the root and ending at an internal node.

# Search Example II


Unsuccessful `TreeSearch(25)`




- An unsuccessful search traverses a path starting at the root and ending at an external node

# Cost of Search

---



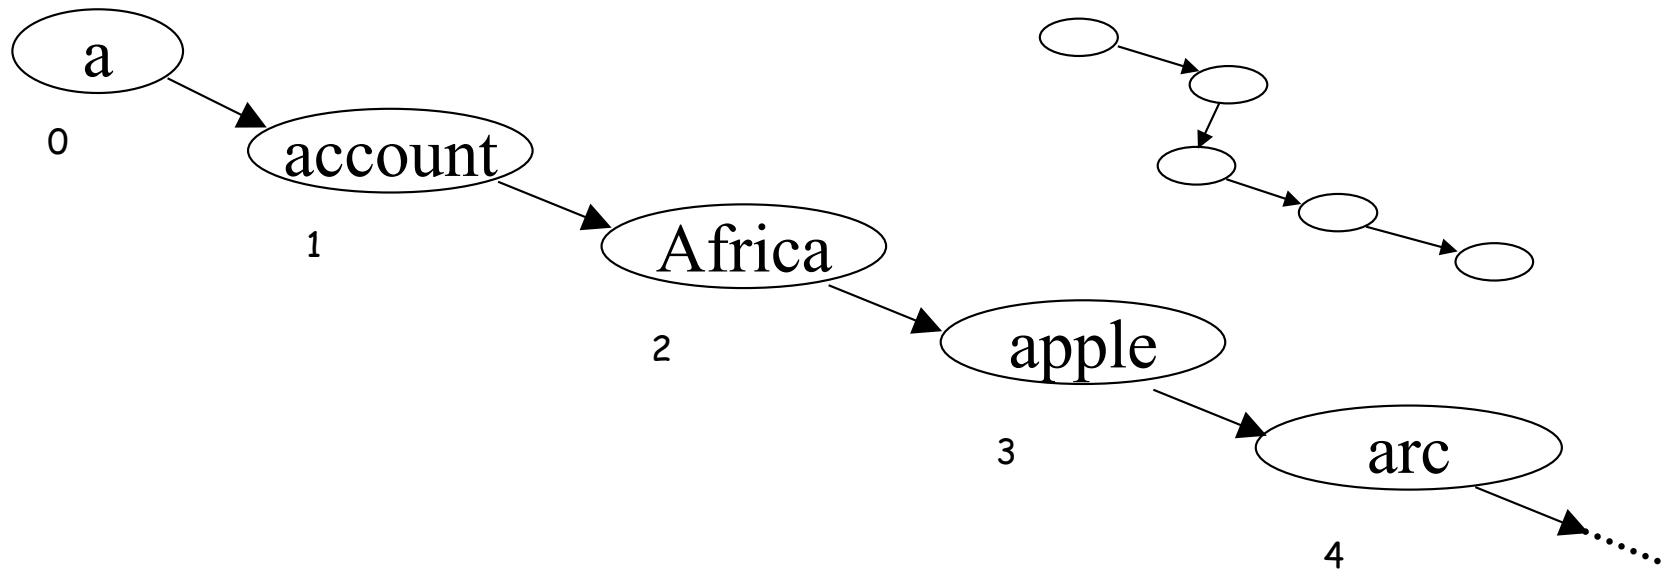
Worst tree



Best tree

# Cost of Search: Worst Case

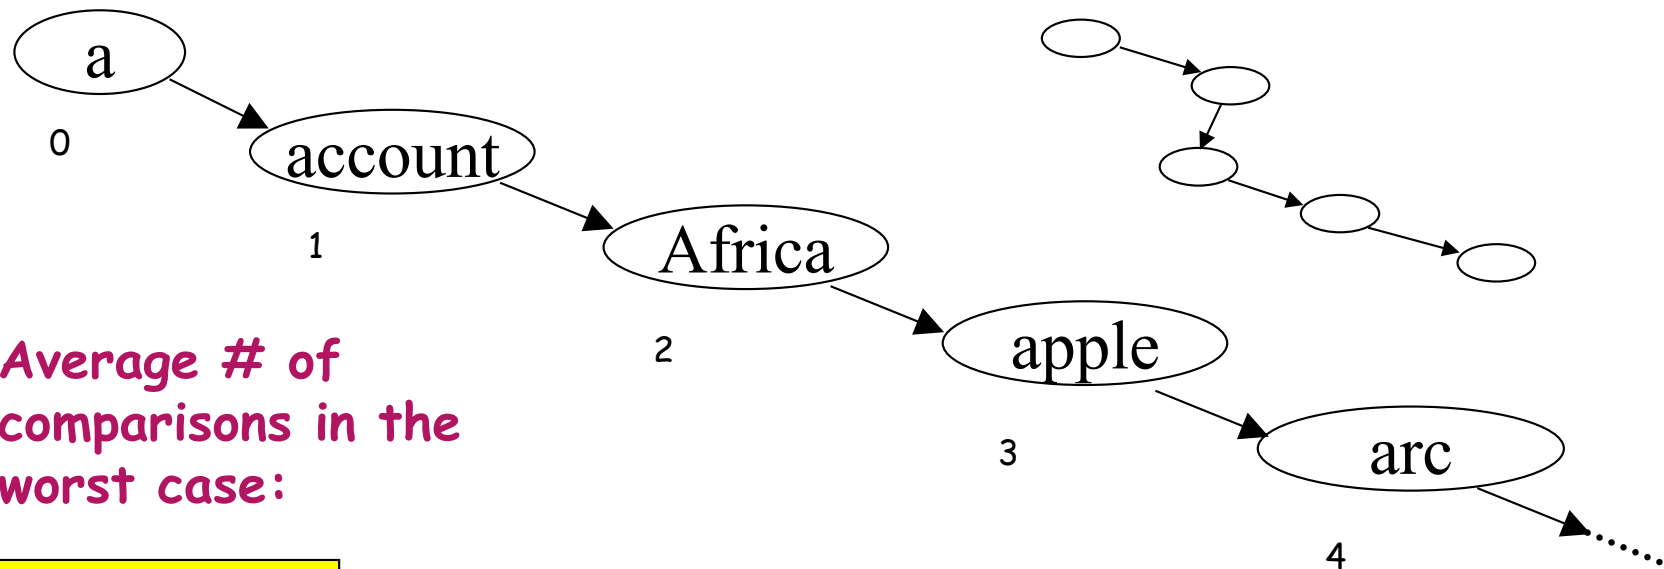
---



Worst possible Tree:  
Worst Case:

$O(n)$

# Cost of Search - Average Worst Case



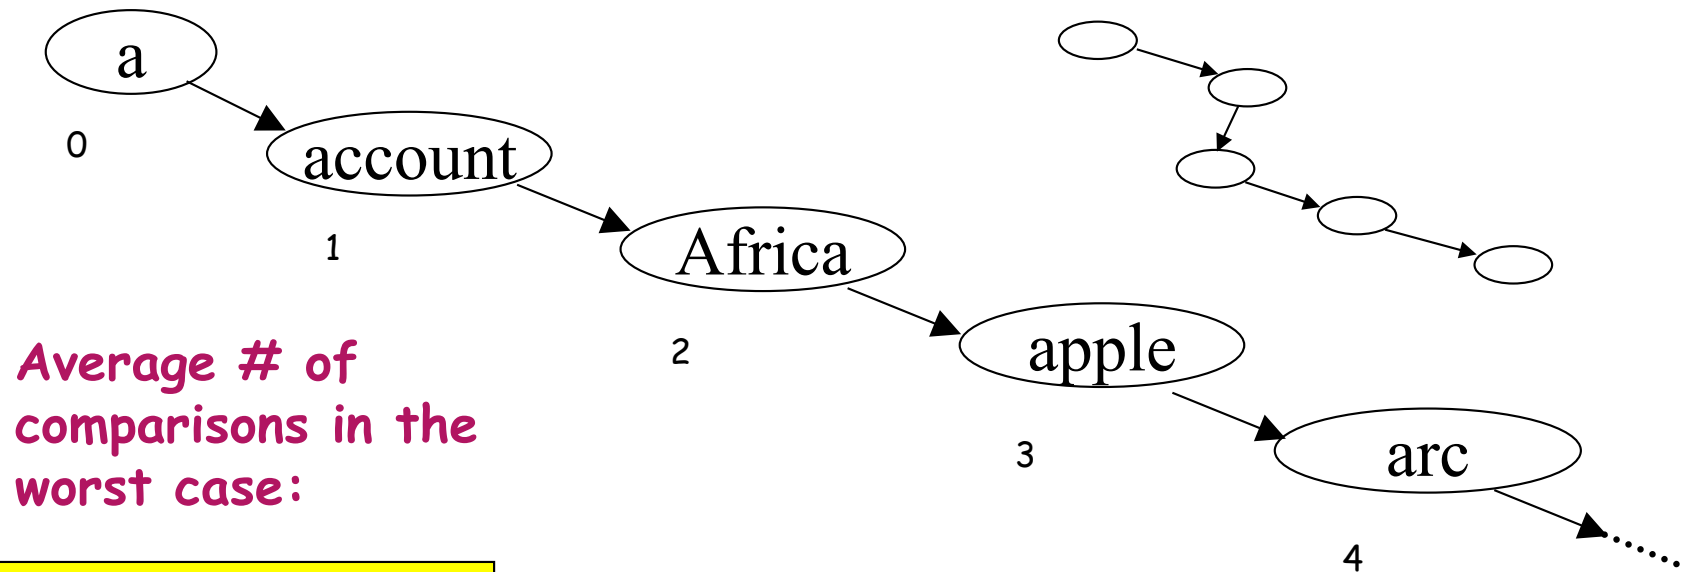
Successful  
search

Path to node  $i$  has length  $i$ , to get there we do  $i$  comparisons

$$\text{Avg cost} = (1/n) \sum i = O(n)$$

# Cost of Search - Average Worst Case

---



Average # of  
comparisons in the  
worst case:


Unsuccessful  
search

An unsuccessful search  
always takes  $O(n)$  comparisons  
for  $n$  internal nodes



# Cost of Search: Best Case

---




Leaves are on the same level or on an adjacent level.

Length of path from root to node  $i = \lfloor \log i \rfloor$

Worst case of the best possible tree:  $O(\log n)$

# Cost of Search: Average Best Case

---



Leaves are on the same level or on an adjacent level.

Length of path from root to node  $i = \lfloor \log i \rfloor$


Comparisons to node  $i$ :  $\log i$

→ Average # of comparisons in the best possible tree

Successful  
search

$$\frac{1}{n} \sum_{i=1}^n \log i = O((n \log n) / n) = O(\log n)$$

Comparisons to node  $i$ : path of length (integer part of)  $\log i$






Ex: to node 2:  $\log 2 = 1$

to node 4,5,6,7:  
 $\log 4 = \log 5 =$   
 $\log 6 = \log 7 = 2$

Comparisons to node  $i$ :  $O(\log i)$

# Cost of Search: Average Best Case

---



Leaves are on the same level or on an adjacent level.

Length of path from root to node  $i = \lfloor \log i \rfloor$

Only paths to external nodes count.

Unsuccessful  
search

always  $O(\log n)$

# Summary

---

Worst tree:

Successful search:

Worst case:  $O(n)$

Average case:  $O(n)$

Unsuccessful search:

Always:  $O(n)$

Best Tree:

Successful search:

Worst case:  $O(\log n)$

Average case:  $O(\log n)$


Unsuccessful search:

Always:  $O(\log n)$

**Note: worst case search for arbitrary binary search tree  $O(n)$**



# Insertion case I

- To perform `TreeInsert(k, v)`, let `w` be the node returned by `TreeSearch(k, T.root())`
- If `w` is external, we know that `k` is not stored in `T`. We call `expandExternal(w, (k, v))` to store `(k, e)` in `w`



# expandExternal(p,(k,v)):

Transform p from an external node into an internal node by creating two new children



expandExternal(p,(k,v)):

if isExternal(p)

create new nodes new1 and new 2

p.left  $\leftarrow$  new1

p.right  $\leftarrow$  new2

store entry (k,v) in p

size  $\leftarrow$  size +2

## Insertion case II

---

- If  $w$  is internal, we know the item with key  $k$  is stored at  $w$ .  
In this case, we just replace the value on this node to the given value  $v$ .

## Insertion in a Binary Search Tree

---

```
Algorithm TreeInsert(k,v)
    p = TreeSearch(root(),k)
    if k == key(p) then
        change p's value to (v)
    else
        expandExternal(p,(k,v)):
```



# Construct a Tree



What would be the result of constructing a tree from repeated insertions of the following sequences?

- a. 5,8,3,7,1,9,2,4,6
- b. 1,2,3,4,5,6,7,8,9
- c. 5,4,6,3,7,2,8,1,9




When do you think trees work best?



# Deletion I

- To perform operation  $\text{remove}(k)$ , we search for key  $k$
- Assume key  $k$  is in the tree, and let  $v$  be the node storing  $k$
- If node  $v$  has a leaf child  $w$ , we remove  $v$  and  $w$  from the tree with operation  $\text{removeAboveExternal}(w)$
- Example: remove 4



removeAboveExternal(v):







removeAboveExternal(v):

```

if isExternal(v) {
  p ← parent(v)
  s ← sibling(v)
  if isRoot(p) {
    s.parent ← null
    root ← s
  }
  else {
    g ← parent(p)
    if (p is leftChild(g)) g.left ← s
    else g.right ← s
    s.parent ← g
  }
  size ← size - 2
}

```






# Deletion II

- We consider the case where the key  $k$  to be removed is stored at a node  $p$  whose children are both internal

- we find the internal node  $w$  that follows  $p$  in an inorder traversal (note  $w$  it does not have a left child!)
- we copy  $entry(w)$  into node  $p$
- we remove node  $w$  and its left child  $z$  (which must be a leaf) by means of operation `removeAboveExternal(z)`

- Example: `remove(3)`

Note: see textbook for different approach: locating node  $w$  that precedes  $p$  in inorder traversal. How would this change the method above?





# Practice, practice, practice...

- a. Delete the 3 from the tree you got in exercise (a) in page 32.
- b. Now delete node 5.



# Cost of Inserting and Deleting = Cost of Search



---

## Summary:

Consider a map with  $n$  items implemented by means of a binary search tree of height  $h$

- the space used is  $O(n)$
- methods `findElement`, `insertItem` and `removeElement` take  $O(h)$  time

The height  $h$  is  $O(n)$  in the worst case and  $O(\log n)$  in the best case




# Conclusion

---


- To achieve good running time, we need to keep the tree *balanced*, i.e., with  $O(\log n)$  height.
- Various balancing schemes can be explored:  
*AVL trees* and *red-black trees* are a balanced binary search trees: their height is  $O(\log n)$   
A *(2,4)-tree* is a search tree (not binary, each internal node has 2, 3 or 4 kids); its height is also  $O(\log n)$ .

Using simply a binary search tree gives worst case running time  $O(n)$  for search, insert and delete operations!


Delete 13



Delete 13



Insert 16



Add 16

